

Data Exploration and Cleaning

```
In [1]: import pandas
```

Daily Activity

```
In [2]: daily_activity = pandas.read_csv("dailyActivity_merged.csv")
```

```
In [3]: daily_activity.head()
```

```
Out[3]:
```

	Id	ActivityDate	TotalSteps	TotalDistance	TrackerDistance	LoggedActivitiesDi
0	1503960366	4/12/2016	13162	8.50	8.50	
1	1503960366	4/13/2016	10735	6.97	6.97	
2	1503960366	4/14/2016	10460	6.74	6.74	
3	1503960366	4/15/2016	9762	6.28	6.28	
4	1503960366	4/16/2016	12669	8.16	8.16	

```
In [4]: daily_activity.dtypes
```

```
Out[4]: Id                int64
ActivityDate            object
TotalSteps              int64
TotalDistance           float64
TrackerDistance         float64
LoggedActivitiesDistance float64
VeryActiveDistance      float64
ModeratelyActiveDistance float64
LightActiveDistance     float64
SedentaryActiveDistance float64
VeryActiveMinutes       int64
FairlyActiveMinutes     int64
LightlyActiveMinutes    int64
SedentaryMinutes        int64
Calories                int64
dtype: object
```

```
In [5]: daily_activity["ActivityDate"] = pandas.to_datetime(daily_activity["ActivityDate"],
```

```
In [6]: daily_activity.dtypes
```

```
Out[6]: Id                                int64
        ActivityDate                      datetime64[ns]
        TotalSteps                        int64
        TotalDistance                     float64
        TrackerDistance                   float64
        LoggedActivitiesDistance          float64
        VeryActiveDistance                 float64
        ModeratelyActiveDistance          float64
        LightActiveDistance               float64
        SedentaryActiveDistance            float64
        VeryActiveMinutes                  int64
        FairlyActiveMinutes                int64
        LightlyActiveMinutes              int64
        SedentaryMinutes                   int64
        Calories                          int64
        dtype: object
```

```
In [7]: daily_activity.shape
```

```
Out[7]: (940, 15)
```

```
In [8]: len(daily_activity["Id"].unique())
```

```
Out[8]: 33
```

```
In [9]: daily_activity["Id"].value_counts()
```

```
Out[9]: 1503960366    31
        4319703577    31
        8583815059    31
        8378563200    31
        8053475328    31
        7086361926    31
        6962181067    31
        5553957443    31
        4702921684    31
        4558609924    31
        1624580081    31
        4388161847    31
        4445114986    31
        8877689391    31
        1927972279    31
        2873212765    31
        2320127002    31
        4020332650    31
        2026352035    31
        1844505072    31
        2022484408    31
        3977333714    30
        1644430081    30
        5577150313    30
        8792009665    29
        6290855005    29
        6117666160    28
        6775888955    26
        7007744171    26
        3372868164    20
        8253242879    19
        2347167796    18
        4057192912     4
        Name: Id, dtype: int64
```

```
In [10]: daily_activity.isna().values.any()
```

```
Out[10]: False
```

```
In [11]: daily_activity.duplicated().any()
```

```
Out[11]: False
```

```
In [12]: daily_activity = daily_activity.loc[:, ("Id", "ActivityDate", "TotalSteps", "Sedent
```

```
In [13]: daily_activity_steps = daily_activity.loc[:, ("Id", "ActivityDate", "TotalSteps")]
        daily_activity_steps["ActivityDate"] = daily_activity_steps["ActivityDate"].dt.date
```

Checking the Consistency of Data on Calories

```
In [14]: daily_activity_calories = daily_activity.copy().loc[:, ("Id", "ActivityDate", "Calo
        daily_activity_calories["ActivityDate"] = daily_activity_calories["ActivityDate"].d
```

```

In [15]: daily_calories = pandas.read_csv("dailyCalories_merged.csv")

In [16]: daily_calories["ActivityDay"] = pandas.to_datetime(daily_calories["ActivityDay"], f

In [17]: daily_calories["ActivityDay"] = daily_calories["ActivityDay"].dt.date
daily_calories = daily_calories.rename(columns={"Calories": "DailyCalories"})

In [18]: hourly_calories = pandas.read_csv("hourlyCalories_merged.csv")

In [19]: hourly_calories["ActivityHour"] = pandas.to_datetime(hourly_calories["ActivityHour"]

In [20]: hourly_calories_added = hourly_calories.copy()
hourly_calories_added["ActivityHour"] = hourly_calories_added["ActivityHour"].dt.da

In [21]: hourly_calories_added = hourly_calories_added.groupby(["Id", "ActivityHour"]).sum()

In [22]: minute_calories_narrow = pandas.read_csv("minuteCaloriesNarrow_merged.csv")

In [23]: minute_calories_narrow["ActivityMinute"] = pandas.to_datetime(minute_calories_narro

In [24]: minute_calories_added = minute_calories_narrow.copy()
minute_calories_added["ActivityMinute"] = minute_calories_added["ActivityMinute"].d

In [25]: minute_calories_added = minute_calories_added.groupby(["Id", "ActivityMinute"]).sum

In [26]: calorie_consistency_test = pandas.merge(daily_activity_calories, daily_calories, le
calorie_consistency_test = pandas.merge(calorie_consistency_test, hourly_calories_a
calorie_consistency_test = pandas.merge(calorie_consistency_test, minute_calories_a
calorie_consistency_test.head()

```

```

Out[26]:

```

	Id	ActivityDate	Calories	ActivityDay	DailyCalories	ActivityHour	HourlyCalo
0	1503960366	2016-04-12	1985	2016-04-12	1985	2016-04-12	
1	1503960366	2016-04-13	1797	2016-04-13	1797	2016-04-13	
2	1503960366	2016-04-14	1776	2016-04-14	1776	2016-04-14	
3	1503960366	2016-04-15	1745	2016-04-15	1745	2016-04-15	
4	1503960366	2016-04-16	1863	2016-04-16	1863	2016-04-16	

```

In [27]: calorie_consistency_test[(calorie_consistency_test["Calories"] == calorie_consisten
(calorie_consistency_test["Calories"] == (calor

```

```
Out[27]:
```

	Id	ActivityDate	Calories	ActivityDay	DailyCalories	ActivityHour	HourlyCal
3	1503960366	2016-04-15	1745	2016-04-15	1745	2016-04-15	
15	1503960366	2016-04-27	2159	2016-04-27	2159	2016-04-27	
36	1624580081	2016-04-18	1604	2016-04-18	1604	2016-04-18	
44	1624580081	2016-04-26	1402	2016-04-26	1402	2016-04-26	
50	1624580081	2016-05-02	1497	2016-05-02	1497	2016-05-02	

```
In [28]: calorie_consistency_test[calorie_consistency_test["Calories"] != calorie_consistenc
```

```
Out[28]:
```

	Id	ActivityDate	Calories	ActivityDay	DailyCalories	ActivityHour	HourlyCalo
0	1503960366	2016-04-12	1985	2016-04-12	1985	2016-04-12	
1	1503960366	2016-04-13	1797	2016-04-13	1797	2016-04-13	
4	1503960366	2016-04-16	1863	2016-04-16	1863	2016-04-16	
5	1503960366	2016-04-17	1728	2016-04-17	1728	2016-04-17	
6	1503960366	2016-04-18	1921	2016-04-18	1921	2016-04-18	

```
In [29]: calorie_consistency_test['MergedToDailyPercentDifference'] = calorie_consistency_te
lambda row: (abs(row['Calories'] - row['DailyCalories']) / row['DailyCalories'])
```

```
In [30]: calorie_consistency_test['MergedtoHourlyPercentDifference'] = calorie_consistency_t
lambda row: (abs(row['Calories'] - row['HourlyCaloriesAdded']) / row['HourlyCal
```

```
In [31]: calorie_consistency_test['MergedtoMinutePercentDifference'] = calorie_consistency_t
lambda row: (abs(row['Calories'] - row['MinuteCaloriesAdded']) / row['MinuteCal
```

```
In [32]: calorie_consistency_test.loc[:, ("MergedToDailyPercentDifference", "MergedtoHourlyP
```

```
Out[32]:
```

	MergedToDailyPercentDifference	MergedtoHourlyPercentDifference	MergedtoMinute
count	934.0	934.000000	
mean	0.0	0.966506	
std	0.0	6.224956	
min	0.0	0.000000	
25%	0.0	0.040469	
50%	0.0	0.087356	
75%	0.0	0.192725	
max	0.0	107.188161	

```
In [33]: print(f'Mismatch: {calorie_consistency_test[calorie_consistency_test["MergedtoHourl
print(f'Total: {calorie_consistency_test["Id"].count()}')
```

Mismatch: 802

Total: 934

765 out of 894 observations contain a mismatch between daily calories, and daily calories calculated from hourly or minute calories. Possible reasons for such discrepancy are rounding error and mistakes during data entry.

```
In [34]: calorie_consistency_test_dropped_high_errors = calorie_consistency_test[(calorie_con  
calorie_consistency_test_dropped_high_errors.loc[:, ("MergedToDailyPercentDifference
```

```
Out[34]:
```

	MergedToDailyPercentDifference	MergedtoHourlyPercentDifference	MergedtoMinute
--	--------------------------------	---------------------------------	----------------

count	905.0	905.000000
mean	0.0	0.186183
std	0.0	0.448649
min	0.0	0.000000
25%	0.0	0.039730
50%	0.0	0.079156
75%	0.0	0.172513
max	0.0	4.532677

```
In [35]: daily_activity = daily_activity[daily_activity.index.isin(calorie_consistency_test_
```

```
In [36]: daily_activity.shape
```

```
Out[36]: (905, 5)
```

Observations with more than 5% difference between daily calories, and calories calculated from hourly or minute calories were removed.

Checking the Consistency of Data on Daily Steps

```
In [37]: daily_steps = pandas.read_csv("dailySteps_merged.csv")
```

```
In [38]: daily_steps["ActivityDay"] = pandas.to_datetime(daily_steps["ActivityDay"], format
```

```
In [39]: daily_steps["ActivityDay"] = daily_steps["ActivityDay"].dt.date  
daily_steps = daily_steps.rename(columns={"StepTotal": "DailySteps"})
```

```
In [40]: hourly_steps_added = pandas.read_csv("hourlySteps_merged.csv")
```

```
In [41]: hourly_steps_added["ActivityHour"] = pandas.to_datetime(hourly_steps_added["Activit
```

```
In [42]: hourly_steps_added["ActivityHour"] = hourly_steps_added["ActivityHour"].dt.date
```

```
In [43]: hourly_steps_added = hourly_steps_added.groupby(["Id", "ActivityHour"]).sum().renam
```

```
In [44]: minute_steps_added = pandas.read_csv("minuteStepsNarrow_merged.csv")
```

```
In [45]: minute_steps_added["ActivityMinute"] = pandas.to_datetime(minute_steps_added["Activ
```

```
In [46]: minute_steps_added["ActivityMinute"] = minute_steps_added["ActivityMinute"].dt.date
```

```
In [47]: minute_steps_added = minute_steps_added.groupby(["Id", "ActivityMinute"]).sum()  
minute_steps_added = minute_steps_added.rename(columns={"Steps": "MinuteStepsAdded"})  
minute_steps_added
```

```
Out[47]:
```

	Id	ActivityMinute	MinuteStepsAdded
--	-----------	-----------------------	-------------------------

0	1503960366	2016-04-12	13158
1	1503960366	2016-04-13	10735
2	1503960366	2016-04-14	10460
3	1503960366	2016-04-15	9685
4	1503960366	2016-04-16	12669
...
929	8877689391	2016-05-08	10665
930	8877689391	2016-05-09	20156
931	8877689391	2016-05-10	10693
932	8877689391	2016-05-11	21391
933	8877689391	2016-05-12	7120

934 rows × 3 columns

```
In [48]: step_consistency_test = pandas.merge(daily_activity_steps, daily_steps, left_on=["I  
step_consistency_test = pandas.merge(step_consistency_test, hourly_steps_added, lef  
step_consistency_test = pandas.merge(step_consistency_test, minute_steps_added, lef  
step_consistency_test.head()
```

```
Out[48]:
```

	Id	ActivityDate	TotalSteps	ActivityDay	DailySteps	ActivityHour	HourlyStep
--	-----------	---------------------	-------------------	--------------------	-------------------	---------------------	-------------------

0	1503960366	2016-04-12	13162	2016-04-12	13162	2016-04-12	
1	1503960366	2016-04-13	10735	2016-04-13	10735	2016-04-13	
2	1503960366	2016-04-14	10460	2016-04-14	10460	2016-04-14	
3	1503960366	2016-04-15	9762	2016-04-15	9762	2016-04-15	
4	1503960366	2016-04-16	12669	2016-04-16	12669	2016-04-16	

```
In [49]: step_consistency_test[(step_consistency_test["TotalSteps"] == step_consistency_test  
                                (step_consistency_test["TotalSteps"] == (step_c
```

Out[49]:

	Id	ActivityDate	TotalSteps	ActivityDay	DailySteps	ActivityHour	HourlyStep
1	1503960366	2016-04-13	10735	2016-04-13	10735	2016-04-13	
2	1503960366	2016-04-14	10460	2016-04-14	10460	2016-04-14	
4	1503960366	2016-04-16	12669	2016-04-16	12669	2016-04-16	
5	1503960366	2016-04-17	9705	2016-04-17	9705	2016-04-17	
6	1503960366	2016-04-18	13019	2016-04-18	13019	2016-04-18	

In [50]: `step_consistency_test[(step_consistency_test["TotalSteps"] != step_consistency_test["TotalSteps"] != (step_c`

Out[50]:

	Id	ActivityDate	TotalSteps	ActivityDay	DailySteps	ActivityHour	HourlyStep
0	1503960366	2016-04-12	13162	2016-04-12	13162	2016-04-12	
3	1503960366	2016-04-15	9762	2016-04-15	9762	2016-04-15	
15	1503960366	2016-04-27	18134	2016-04-27	18134	2016-04-27	
16	1503960366	2016-04-28	13154	2016-04-28	13154	2016-04-28	
20	1503960366	2016-05-02	14727	2016-05-02	14727	2016-05-02	

In [51]: `step_consistency_test[(step_consistency_test["TotalSteps"] > 0) & (step_consistency`

Out[51]:

	Id	ActivityDate	TotalSteps	ActivityDay	DailySteps	ActivityHour	HourlyStep
379	4319703577	2016-04-12	7753	2016-04-12	7753	2016-04-12	
410	4388161847	2016-04-12	10122	2016-04-12	10122	2016-04-12	
849	8583815059	2016-04-16	5319	2016-04-16	5319	2016-04-16	
850	8583815059	2016-04-17	3008	2016-04-17	3008	2016-04-17	
865	8583815059	2016-05-02	8469	2016-05-02	8469	2016-05-02	

Calculating the % difference between daily steps and daily steps calculated from hourly and minute steps. If daily step count was >0, but daily step count calculated from hourly or minute steps was 0,value of 6% was placed for that row as observations with higher than 5% difference will be filtered out for further analysis.

In [52]: `step_consistency_test['MergedToDailyPercentDifference'] = step_consistency_test.app
lambda row: (abs(row['TotalSteps'] - row['DailySteps']) / row['DailySteps']) *`

In [53]: `step_consistency_test['MergedtoHourlyPercentDifference'] = step_consistency_test.ap
lambda row: (abs(row['TotalSteps'] - row['HourlyStepsAdded']) / row['HourlyStep`

In [54]: `step_consistency_test['MergedtoMinutePercentDifference'] = step_consistency_test.ap
lambda row: (abs(row['TotalSteps'] - row['MinuteStepsAdded']) / row['MinuteStep`


```
In [55]: step_consistency_test.head()
```

```
Out[55]:
```

	Id	ActivityDate	TotalSteps	ActivityDay	DailySteps	ActivityHour	HourlySteps
0	1503960366	2016-04-12	13162	2016-04-12	13162	2016-04-12	
1	1503960366	2016-04-13	10735	2016-04-13	10735	2016-04-13	
2	1503960366	2016-04-14	10460	2016-04-14	10460	2016-04-14	
3	1503960366	2016-04-15	9762	2016-04-15	9762	2016-04-15	
4	1503960366	2016-04-16	12669	2016-04-16	12669	2016-04-16	

```
In [56]: step_consistency_test.loc[:, ("MergedToDailyPercentDifference", "MergedtoHourlyPerce
```

```
Out[56]:
```

	MergedToDailyPercentDifference	MergedtoHourlyPercentDifference	MergedtoMinuteDifference
count	934.0	934.000000	
mean	0.0	3.176913	
std	0.0	34.616543	
min	0.0	0.000000	
25%	0.0	0.000000	
50%	0.0	0.000000	
75%	0.0	0.000000	
max	0.0	724.618447	

```
In [57]: print(f'Mismatch: {step_consistency_test[step_consistency_test["MergedtoHourlyPerce
print(f'Total: {step_consistency_test["Id"].count()}')
```

Mismatch: 159

Total: 934

159 out of 934 observations contain a mismatch between daily steps, and daily steps calculated from hourly or minute steps. Possible reasons for such discrepancy are rounding error and mistakes during data entry.

```
In [58]: step_consistency_test_dropped_high_errors = step_consistency_test[(step_consistency_
```

```
In [59]: step_consistency_test_dropped_high_errors.loc[:, ("MergedToDailyPercentDifference",
```

Out[59]:	MergedToDailyPercentDifference	MergedtoHourlyPercentDifference	MergedtoMinute
count	900.0	900.000000	
mean	0.0	0.081552	
std	0.0	0.354611	
min	0.0	0.000000	
25%	0.0	0.000000	
50%	0.0	0.000000	
75%	0.0	0.000000	
max	0.0	4.767442	

```
In [60]: daily_activity = daily_activity[daily_activity.index.isin(step_consistency_test_dro
```

```
In [61]: len(daily_activity["Id"].unique())
```

```
Out[61]: 33
```

```
In [62]: (daily_activity["Id"].value_counts() >= 28).sum()
```

```
Out[62]: 23
```

Observations with more than 5% difference between daily steps, and daily steps calculated from hourly steps were removed.

Hourly Intensities

```
In [63]: hourly_intensities = pandas.read_csv("hourlyIntensities_merged.csv")
```

```
In [64]: hourly_intensities.head()
```

Out[64]:		Id	ActivityHour	TotalIntensity	AverageIntensity
	0	1503960366	4/12/2016 12:00:00 AM	20	0.333333
	1	1503960366	4/12/2016 1:00:00 AM	8	0.133333
	2	1503960366	4/12/2016 2:00:00 AM	7	0.116667
	3	1503960366	4/12/2016 3:00:00 AM	0	0.000000
	4	1503960366	4/12/2016 4:00:00 AM	0	0.000000

```
In [65]: hourly_intensities.dtypes
```

```
Out[65]: Id                int64
ActivityHour             object
TotalIntensity           int64
AverageIntensity         float64
dtype: object
```

```
In [66]: hourly_intensities["ActivityHour"] = pandas.to_datetime(hourly_intensities["Activit
```

```
In [67]: hourly_intensities.dtypes
```

```
Out[67]: Id                int64
ActivityHour             datetime64[ns]
TotalIntensity           int64
AverageIntensity         float64
dtype: object
```

```
In [68]: hourly_intensities.shape
```

```
Out[68]: (22099, 4)
```

```
In [69]: len(hourly_intensities["Id"].unique())
```

```
Out[69]: 33
```

```
In [70]: hourly_intensities.isna().values.any()
```

```
Out[70]: False
```

```
In [71]: hourly_intensities.duplicated().any()
```

```
Out[71]: False
```

Minute MET

```
In [72]: minute_met = pandas.read_csv("minuteMETsNarrow_merged.csv")
```

```
In [73]: minute_met.head()
```

```
Out[73]:
```

	Id	ActivityMinute	METs
0	1503960366	4/12/2016 12:00:00 AM	10
1	1503960366	4/12/2016 12:01:00 AM	10
2	1503960366	4/12/2016 12:02:00 AM	10
3	1503960366	4/12/2016 12:03:00 AM	10
4	1503960366	4/12/2016 12:04:00 AM	10

```
In [74]: minute_met.dtypes
```

```
Out[74]: Id                int64
ActivityMinute          object
METs                   int64
dtype: object
```

```
In [75]: minute_met["ActivityMinute"] = pandas.to_datetime(minute_met["ActivityMinute"], for
```

```
In [76]: minute_met.dtypes
```

```
Out[76]: Id                int64
ActivityMinute    datetime64[ns]
METs              int64
dtype: object
```

```
In [77]: minute_met.shape
```

```
Out[77]: (1325580, 3)
```

```
In [78]: len(minute_met["Id"].unique())
```

```
Out[78]: 33
```

```
In [79]: minute_met.isna().values.any()
```

```
Out[79]: False
```

```
In [80]: minute_met.duplicated().any()
```

```
Out[80]: False
```

```
In [81]: minute_met_min = minute_met.copy()
```

```
In [82]: minute_met_min["ActivityMinute"] = minute_met_min["ActivityMinute"].dt.strftime("%H
```

```
In [83]: grouped_met_min = minute_met_min.groupby(["ActivityMinute"])["METs"].describe()
```

```
In [84]: grouped_met_min[grouped_met_min.index.isin(["00:00:00", "02:00:00", "04:00:00", "6:
```

The history saving thread hit an unexpected error (OperationalError('database is locked')).History will not be written to the database.

Out[84]:

	count	mean	std	min	25%	50%	75%	max
--	-------	------	-----	-----	-----	-----	-----	-----

ActivityMinute

00:00:00	934.0	11.688437	9.057299	10.0	10.0	10.0	10.0	133.0
02:00:00	933.0	10.534834	3.888483	10.0	10.0	10.0	10.0	80.0
04:00:00	932.0	10.211373	2.178737	10.0	10.0	10.0	10.0	52.0
08:00:00	931.0	15.264232	11.634348	10.0	10.0	10.0	12.0	96.0
10:00:00	929.0	16.525296	13.646067	10.0	10.0	10.0	13.0	138.0
12:00:00	922.0	17.200651	13.476641	10.0	10.0	11.0	24.0	113.0
14:00:00	918.0	18.245098	15.911281	10.0	10.0	11.0	24.0	129.0
16:00:00	907.0	16.471885	12.949641	10.0	10.0	11.0	13.0	129.0
18:00:00	906.0	18.584989	15.690021	10.0	10.0	11.0	24.0	133.0
20:00:00	906.0	16.835541	13.277177	10.0	10.0	11.0	14.0	125.0
22:00:00	904.0	14.613938	12.037606	10.0	10.0	10.0	12.0	140.0

MET is defined as a rate of energy expended during an activity compared to the rate of energy expended during rest. The summary table above shows that during sleep hours (~midnight to 8 am) there is no difference in MET values between the quartiles, which is expected as energy expenditure of sleeping people is virtually the same. An increase in the MET values is seen in the 50th and 75th percentiles in later hours, which is likely due to the physical activity of more active people. Although these observations are consistent with the context, at rest and sleep, MET is usually around 1, not 10, suggesting that there is an error in data, or a different definition/equation for MET was used in this dataset. As metadata for this dataset not available, an assumption was made that MET values in the dataset were multiplied by 10. To correct this for further analysis, MET values were divided by 10.

```
In [85]: minute_met_h = minute_met.copy()
```

```
In [86]: minute_met_h["METs"] = minute_met_h["METs"] / 10
```

```
In [87]: minute_met_h["ActivityMinute"] = minute_met_h["ActivityMinute"].dt.strftime("%Y-%m-
```

```
In [88]: minute_met_h = minute_met_h.rename(columns={"ActivityMinute": "Date"})
```

```
In [89]: minute_met_h.dtypes
```

```
Out[89]: Id          int64
Date          object
METs         float64
dtype: object
```

```
In [90]: average_daily_met = minute_met_h.groupby(["Id", "Date"]).mean()
```

```
In [91]: average_daily_met = average_daily_met.reset_index()
```

```
In [92]: average_daily_met["Date"] = pandas.to_datetime(average_daily_met["Date"])
```

Daily Sleep

```
In [93]: daily_sleep = pandas.read_csv("sleepDay_merged.csv")
```

```
In [94]: daily_sleep.head()
```

```
Out[94]:
```

	Id	SleepDay	TotalSleepRecords	TotalMinutesAsleep	TotalTimeInBed
0	1503960366	4/12/2016 12:00:00 AM	1	327	346
1	1503960366	4/13/2016 12:00:00 AM	2	384	407
2	1503960366	4/15/2016 12:00:00 AM	1	412	442
3	1503960366	4/16/2016 12:00:00 AM	2	340	367
4	1503960366	4/17/2016 12:00:00 AM	1	700	712

```
In [95]: daily_sleep.dtypes
```

```
Out[95]: Id                int64
SleepDay                object
TotalSleepRecords       int64
TotalMinutesAsleep       int64
TotalTimeInBed          int64
dtype: object
```

```
In [96]: daily_sleep["SleepDay"] = pandas.to_datetime(daily_sleep["SleepDay"], format = "%m/
```

```
In [97]: daily_sleep.dtypes
```

```
Out[97]: Id                int64
SleepDay                datetime64[ns]
TotalSleepRecords       int64
TotalMinutesAsleep       int64
TotalTimeInBed          int64
dtype: object
```

```
In [98]: daily_sleep.shape
```

```
Out[98]: (413, 5)
```

```
In [99]: len(daily_sleep["Id"].unique())
```

Out[99]: 24

```
In [100]: daily_sleep["Id"].value_counts()
```

```
Out[100]: 8378563200    32
          6962181067    31
          5553957443    31
          4702921684    28
          2026352035    28
          3977333714    28
          4445114986    28
          5577150313    26
          4319703577    26
          1503960366    25
          7086361926    24
          4388161847    24
          6117666160    18
          8792009665    15
          2347167796    15
          4020332650     8
          1927972279     5
          4558609924     5
          1644430081     4
          6775888955     3
          8053475328     3
          1844505072     3
          7007744171     2
          2320127002     1
          Name: Id, dtype: int64
```

```
In [101]: daily_sleep.isna().values.any()
```

Out[101]: False

```
In [102]: daily_sleep.duplicated().any()
```

Out[102]: True

```
In [103]: daily_sleep[daily_sleep.duplicated(keep=False)]
```

```
Out[103]:
```

	Id	SleepDay	TotalSleepRecords	TotalMinutesAsleep	TotalTimeInBed
160	4388161847	2016-05-05	1	471	495
161	4388161847	2016-05-05	1	471	495
222	4702921684	2016-05-07	1	520	543
223	4702921684	2016-05-07	1	520	543
379	8378563200	2016-04-25	1	388	402
380	8378563200	2016-04-25	1	388	402

```
In [104]: daily_sleep = daily_sleep.drop_duplicates()
```

```
In [105... daily_sleep.duplicated().any()
```

```
Out[105]: False
```

Weight Log

```
In [106... weight_log = pandas.read_csv("weightLogInfo_merged.csv")
```

```
In [107... weight_log.head()
```

```
Out[107]:
```

	Id	Date	WeightKg	WeightPounds	Fat	BMI	IsManualReport		
0	1503960366	5/2/2016 11:59:59 PM	52.599998	115.963147	22.0	22.650000	True	1	
1	1503960366	5/3/2016 11:59:59 PM	52.599998	115.963147	NaN	22.650000	True	1	
2	1927972279	4/13/2016 1:08:52 AM	133.500000	294.317120	NaN	47.540001	False	1	
3	2873212765	4/21/2016 11:59:59 PM	56.700001	125.002104	NaN	21.450001	True	1	
4	2873212765	5/12/2016 11:59:59 PM	57.299999	126.324875	NaN	21.690001	True	1	

```
In [108... weight_log.dtypes
```

```
Out[108]: Id                int64
Date                object
WeightKg            float64
WeightPounds        float64
Fat                 float64
BMI                 float64
IsManualReport      bool
LogId               int64
dtype: object
```

```
In [109... weight_log["Date"] = pandas.to_datetime(weight_log["Date"], format = "%m/%d/%Y %I:%M:%S")
```

```
In [110... weight_log.dtypes
```



```
Out[110]: Id                int64
          Date              datetime64[ns]
          WeightKg          float64
          WeightPounds      float64
          Fat               float64
          BMI               float64
          IsManualReport    bool
          LogId             int64
          dtype: object
```

```
In [111]: weight_log.shape
```

```
Out[111]: (67, 8)
```

```
In [112]: len(weight_log["Id"].unique())
```

```
Out[112]: 8
```

```
In [113]: weight_log["Id"].value_counts()
```

```
Out[113]: 6962181067    30
          8877689391    24
          4558609924     5
          1503960366     2
          2873212765     2
          4319703577     2
          1927972279     1
          5577150313     1
          Name: Id, dtype: int64
```

```
In [114]: weight_log.isna().any()
```

```
Out[114]: Id                False
          Date              False
          WeightKg          False
          WeightPounds      False
          Fat               True
          BMI               False
          IsManualReport    False
          LogId             False
          dtype: bool
```

```
In [115]: weight_log[weight_log.isna().any(axis=1)].head()
```

Out[115]:

		Id	Date	WeightKg	WeightPounds	Fat	BMI	IsManualReport	
1	1503960366		2016-05-03 23:59:59	52.599998	115.963147	NaN	22.650000	True	146
2	1927972279		2016-04-13 01:08:52	133.500000	294.317120	NaN	47.540001	False	146
3	2873212765		2016-04-21 23:59:59	56.700001	125.002104	NaN	21.450001	True	146
4	2873212765		2016-05-12 23:59:59	57.299999	126.324875	NaN	21.690001	True	146
6	4319703577		2016-05-04 23:59:59	72.300003	159.394222	NaN	27.379999	True	146

```
In [116... weight_log[~weight_log.isna().any(axis=1)]
```

Out[116]:

		Id	Date	WeightKg	WeightPounds	Fat	BMI	IsManualReport	
0	1503960366		2016-05-02 23:59:59	52.599998	115.963147	22.0	22.650000	True	1462
5	4319703577		2016-04-17 23:59:59	72.400002	159.614681	25.0	27.450001	True	1460

```
In [117... weight_log.duplicated().any()
```

Out[117]: False

Exploratory Graphs

```
In [118... import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import f_oneway, ttest_ind
```

Daily Activity and Sleep

Merging Tables

```
In [119... daily_activity_sleep = pandas.merge(daily_sleep, daily_activity, left_on=["Id", "Sl
```

```
In [120... daily_activity_sleep = pandas.merge(daily_activity_sleep, average_daily_met, left_o
```

```
In [121... daily_activity_sleep = daily_activity_sleep.drop(["ActivityDate", "Date"], axis=1)
```

```
In [122... daily_activity_sleep.rename(columns={"SleepDay":"Date"}, inplace=True)
```

```
In [123... daily_activity_sleep = daily_activity_sleep.loc[:, ("Id", "TotalMinutesAsleep", "To  
daily_activity_sleep.describe()
```

```
Out[123]:
```

	Id	TotalMinutesAsleep	TotalSteps	SedentaryMinutes	Calories
count	3.930000e+02	393.000000	393.000000	393.000000	393.000000
mean	5.003939e+09	418.694656	8545.432570	716.277354	2387.908397
std	2.066512e+09	119.736737	4178.541642	163.129712	757.246898
min	1.503960e+09	58.000000	42.000000	2.000000	403.000000
25%	3.977334e+09	361.000000	5183.000000	637.000000	1837.000000
50%	4.702922e+09	432.000000	8954.000000	720.000000	2196.000000
75%	6.962181e+09	489.000000	11393.000000	787.000000	2908.000000
max	8.792010e+09	796.000000	22770.000000	1265.000000	4900.000000

```
In [124... len(daily_activity_sleep["Id"].unique())
```

```
Out[124]: 24
```

```
In [125... daily_activity_sleep["Id"].value_counts()
```

```
Out[125]: 6962181067    31
          5553957443    31
          8378563200    29
          3977333714    28
          4702921684    27
          2026352035    27
          4445114986    27
          5577150313    25
          1503960366    24
          7086361926    23
          4319703577    23
          4388161847    22
          6117666160    16
          8792009665    15
          2347167796    14
          4020332650     5
          1927972279     5
          4558609924     5
          1644430081     4
          6775888955     3
          8053475328     3
          1844505072     3
          7007744171     2
          2320127002     1
          Name: Id, dtype: int64
```

```
In [126]: (daily_activity_sleep["Id"].value_counts() >= 14).sum()
```

```
Out[126]: 15
```

```
In [127]: (daily_activity_sleep["Id"].value_counts() >= 21).sum()
```

```
Out[127]: 12
```

```
In [128]: (daily_activity_sleep["Id"].value_counts() >= 28).sum()
```

```
Out[128]: 4
```

Graphs

```
In [129]: x = daily_activity_sleep["SedentaryMinutes"]
          y = daily_activity_sleep["TotalMinutesAsleep"]

          plt.scatter(x, y)


          coefficients = np.polyfit(x, y, deg=1)
          slope = coefficients[0]
          intercept = coefficients[1]
          regression_line = slope * x + intercept

          plt.plot(x, regression_line, color='black')

          r_value = np.corrcoef(x, y)[0, 1]
          plt.text(0.2, 0.1, f'r = {r_value:.2f}', ha='center', va='center', transform=plt.gc
```

```
plt.xlabel("Total Daily Sedentary Minutes")
plt.ylabel("Total Daily Minutes Asleep")

plt.show()
```



There is a negative correlation between sedentary time and sleep duration, suggesting prolonged sitting negatively affects the duration of sleep.

```
In [130]: daily_activity_sleep.head()
```

```
Out[130]:
```

	Id	TotalMinutesAsleep	TotalSteps	SedentaryMinutes	Calories	METs
0	1503960366	327	13162	728	1985	1.752847
1	1503960366	384	10735	776	1797	1.587431
2	1503960366	412	9762	726	1745	1.540972
3	1503960366	340	12669	773	1863	1.645417
4	1503960366	700	9705	539	1728	1.525833

```
In [131]: x = daily_activity_sleep["METs"]
y = daily_activity_sleep["TotalMinutesAsleep"]

plt.scatter(x, y)

coefficients = np.polyfit(x, y, deg=1)
slope = coefficients[0]
```

```

intercept = coefficients[1]
regression_line = slope * x + intercept


plt.plot(x, regression_line, color='black')

r_value = np.corrcoef(x, y)[0, 1]
plt.text(0.8, 0.1, f'r = {r_value:.2f}', ha='center', va='center', transform=plt.gc

plt.xlabel("Daily Average METs/minute")
plt.ylabel("Total Daily Minutes Asleep")

plt.show()

```



To explore the relationship between physical activity and sleep duration, Metabolic Equivalent Tasks (METs) were used as a physical activity metric. MET is defined as a rate of energy expended during an activity compared to the rate of energy expended during rest. 1 REM is equivalent to a person resting in a sedentary state, while activities that require MET > 1.5 are generally considered light physical activity (<http://dx.doi.org/10.2196/36181>).

No obvious relationship is observed between average daily activity expressed in average daily METs and sleep duration. This is in contrast to observations in sedentary time vs sleep duration graph. It is possible that daily sedentary time might not reflect the overall person's physical activity very accurately.

```

In [132... x = daily_activity_sleep["SedentaryMinutes"]
y = daily_activity_sleep["METs"]

```

```
plt.scatter(x, y)


coefficients = np.polyfit(x, y, deg=1)
slope = coefficients[0]
intercept = coefficients[1]
regression_line = slope * x + intercept

plt.plot(x, regression_line, color='black')

r_value = np.corrcoef(x, y)[0, 1]
plt.text(0.8, 0.1, f'r = {r_value:.2f}', ha='center', va='center', transform=plt.gca().transData)

plt.xlabel("Total Daily Sedentary Minutes")
plt.ylabel("Daily Average METs/Minute")

plt.show()
```



```
In [133... x = daily_activity_sleep["SedentaryMinutes"]
y = daily_activity_sleep["TotalSteps"]

plt.scatter(x, y)

coefficients = np.polyfit(x, y, deg=1)
slope = coefficients[0]
intercept = coefficients[1]
regression_line = slope * x + intercept


plt.plot(x, regression_line, color='black')

r_value = np.corrcoef(x, y)[0, 1]
```

```
plt.text(0.8, 0.1, f'r = {r_value:.2f}', ha='center', va='center', transform=plt.gc

plt.xlabel("Total Daily Sedentary Minutes")
plt.ylabel("Total Daily Steps")

plt.show()
```



```
In [134... x = daily_activity_sleep["SedentaryMinutes"]
y = daily_activity_sleep["Calories"]

plt.scatter(x, y)

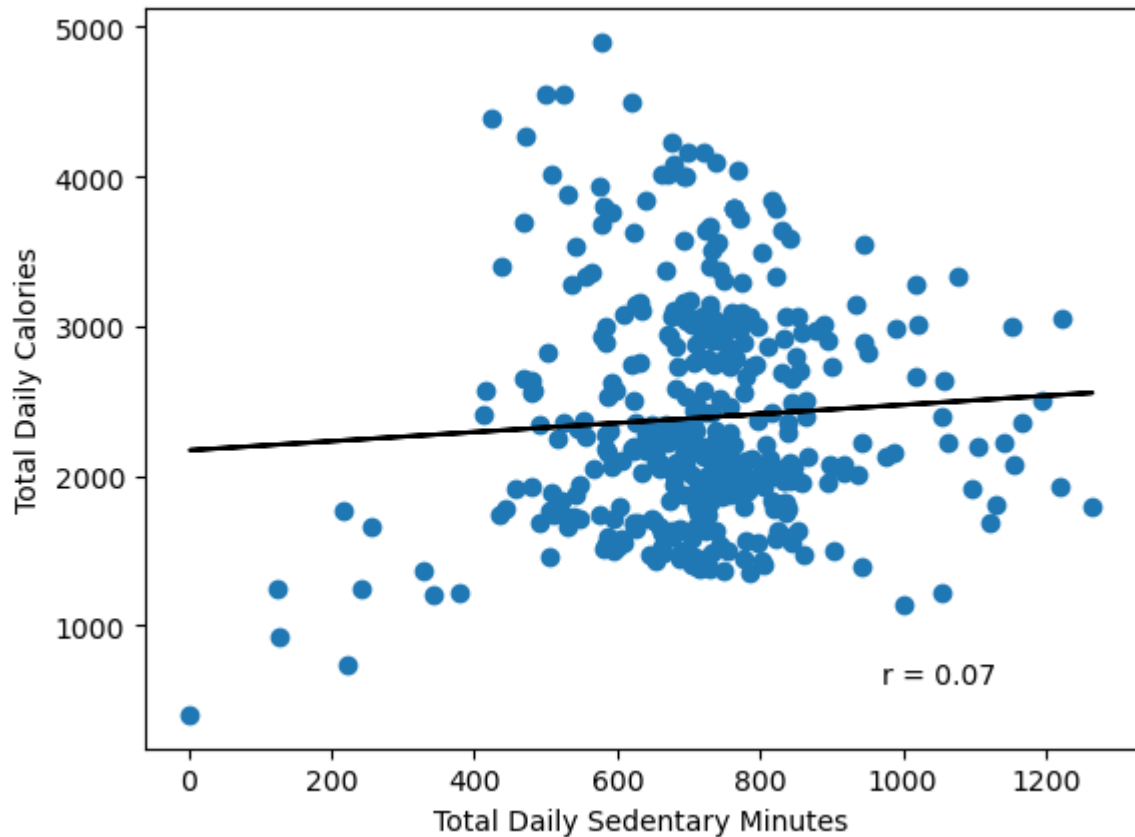
coefficients = np.polyfit(x, y, deg=1)
slope = coefficients[0]
intercept = coefficients[1]
regression_line = slope * x + intercept

plt.plot(x, regression_line, color='black')

r_value = np.corrcoef(x, y)[0, 1]
plt.text(0.8, 0.1, f'r = {r_value:.2f}', ha='center', va='center', transform=plt.gc

plt.xlabel("Total Daily Sedentary Minutes")
plt.ylabel("Total Daily Calories")


plt.show()
```

Indeed, it appears that total daily sedentary time is not a great predictor of persons total daily activity (measured as daily average METs/minute or total daily steps) or total daily energy expenditure (total daily calories).

```
In [135... daily_activity_sleep['METs_Quartile'] = pandas.qcut(daily_activity_sleep['METs'], q
for quartile in ['Q1', 'Q2', 'Q3', 'Q4']:
    plt.scatter(daily_activity_sleep.loc[daily_activity_sleep['METs_Quartile'] == q
                daily_activity_sleep.loc[daily_activity_sleep['METs_Quartile'] == q

plt.xlabel("Daily average METs/Minute")
plt.ylabel("Total Daily Minutes Asleep")
plt.show()
```



Although no upward or downward trends are observed in the average daily MET minutes vs total minutes asleep graph, when data is split into quartiles, data variability seems to decrease as average daily MET minutes increase.

```
In [136... daily_activity_sleep.groupby("METs_Quartile")["TotalMinutesAsleep"].describe()
```

```
Out[136]:
```

	count	mean	std	min	25%	50%	75%	max
METs_Quartile								
Q1	99.0	457.969697	146.282607	59.0	407.5	463.0	526.00	796.0
Q2	99.0	418.232323	111.491267	58.0	358.0	432.0	495.50	700.0
Q3	98.0	396.102041	100.706114	62.0	338.0	417.0	466.75	611.0
Q4	97.0	401.907216	106.604319	74.0	361.0	418.0	472.00	600.0

```
In [137... daily_activity_sleep.groupby("METs_Quartile")["METs"].describe()
```

Out[137]:		count	mean	std	min	25%	50%	75%	max
	METs_Quartile								
	Q1	99.0	1.246579	0.091446	1.018056	1.189826	1.276111	1.321001	1.362014
	Q2	99.0	1.468878	0.053821	1.363611	1.430764	1.478125	1.515382	1.546875
	Q3	98.0	1.612353	0.040929	1.549722	1.577535	1.612500	1.644080	1.685556
	Q4	97.0	1.854094	0.187586	1.688542	1.732083	1.777569	1.905486	2.501667

```
In [138... from scipy.stats import levene
for quartile_1 in ['Q1', 'Q2', 'Q3', 'Q4']:
    for quartile_2 in ['Q1', 'Q2', 'Q3', 'Q4']:
        if quartile_1 != quartile_2:
            statistic, p_value = levene(daily_activity_sleep.loc[daily_activity_sle
            print(f"P-value ({quartile_1} and {quartile_2}): {p_value:.2f}")
```

```
P-value (Q1 and Q2): 0.14
P-value (Q1 and Q3): 0.05
P-value (Q1 and Q4): 0.06
P-value (Q2 and Q1): 0.14
P-value (Q2 and Q3): 0.60
P-value (Q2 and Q4): 0.63
P-value (Q3 and Q1): 0.05
P-value (Q3 and Q2): 0.60
P-value (Q3 and Q4): 0.99
P-value (Q4 and Q1): 0.06
P-value (Q4 and Q2): 0.63
P-value (Q4 and Q3): 0.99
```

Although the average sleep duration is close to the recommended 7 hours, it seems that the standard deviation of data in the first quartile is higher than that in the other quartiles. This suggests, that higher average daily MET minutes might be associated with more consistent nightly sleep duration. This might be important as both under- and oversleeping are known to be detrimental to health.

```
In [139... daily_activity_sleep['SleepQuality'] = ''
daily_activity_sleep.loc[daily_activity_sleep['TotalMinutesAsleep'] > 480, 'SleepQu
daily_activity_sleep.loc[daily_activity_sleep['TotalMinutesAsleep'] < 360, 'SleepQu
daily_activity_sleep.loc[(daily_activity_sleep['TotalMinutesAsleep'] >= 360) & (dai
```

```
In [140... daily_activity_sleep_grouped = daily_activity_sleep.groupby(["METs_Quartile", "Slee
total = daily_activity_sleep_grouped["normal"] + daily_activity_sleep_grouped["over
daily_activity_sleep_grouped = daily_activity_sleep_grouped.div(total, axis=0) * 10
```

```
In [141... x_labels = ['Sedentary', 'LightlyActive', 'FairlyActive', 'VeryActive']
bar_width = 0.2

pos_undersleeping = list(range(len(x_labels)))
pos_normal = [pos + bar_width for pos in pos_undersleeping]
pos_oversleeping = [pos + bar_width for pos in pos_normal]
```

```

plt.bar(pos_undersleeping, daily_activity_sleep_grouped['undersleeping'], width=bar_width, label='undersleeping')
plt.bar(pos_normal, daily_activity_sleep_grouped['normal'], width=bar_width, label='normal')
plt.bar(pos_oversleeping, daily_activity_sleep_grouped['oversleeping'], width=bar_width, label='oversleeping')

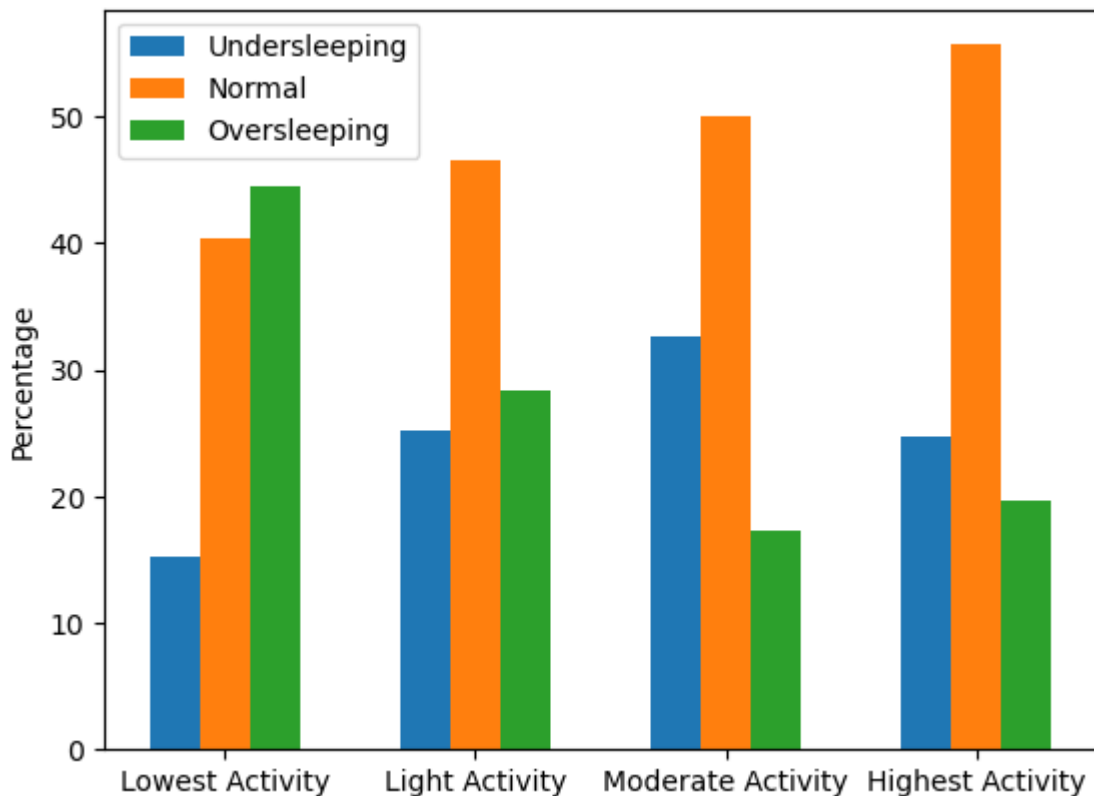
plt.ylabel('Percentage')

plt.xticks([pos + bar_width for pos in range(len(x_labels))], ["Lowest Activity", "Light Activity", "Moderate Activity", "Highest Activity"])

plt.legend()

plt.show()

```



Testing if the increase in the proportion of normal duration sleep events in the highest activity quartile is significantly different from the lowest activity quartile

```

In [142...] daily_activity_sleep_stats = daily_activity_sleep.copy()
daily_activity_sleep_stats['SleepQuality'] = ''
daily_activity_sleep_stats.loc[daily_activity_sleep_stats['TotalMinutesAsleep'] > 4, 'SleepQuality'] = 'good'
daily_activity_sleep_stats.loc[daily_activity_sleep_stats['TotalMinutesAsleep'] < 3, 'SleepQuality'] = 'bad'
daily_activity_sleep_stats.loc[(daily_activity_sleep_stats['TotalMinutesAsleep'] >= 3 & daily_activity_sleep_stats['TotalMinutesAsleep'] <= 4), 'SleepQuality'] = 'normal'

```

Normal duration sleep events reclassified as "good", undersleeping and oversleeping events classified as "bad".

```

In [143...] contingency_table = daily_activity_sleep_stats[["METs_Quartile", "SleepQuality"]].groupby("METs_Quartile").value_counts().unstack()
contingency_table = contingency_table.reset_index()
contingency_table = contingency_table.rename(columns={0:"counts"})
contingency_table = contingency_table.pivot(index='METs_Quartile', columns='SleepQuality', values='counts')

```

```
contingency_table = contingency_table.reset_index()
contingency_table
```

Out[143]: **SleepQuality** **METs_Quartile** **abnormal** **normal**

0	Q1	59	40
1	Q2	53	46
2	Q3	49	49
3	Q4	43	54

```
In [144... import numpy as np
from scipy.stats import chi2_contingency
for quartile_1 in ['Q1', 'Q2', 'Q3', 'Q4']:
    for quartile_2 in ['Q1', 'Q2', 'Q3', 'Q4']:
        if quartile_1 != quartile_2:
            observed_data = contingency_table[contingency_table["METs_Quartile"].is
            chi2, p_value, dof, expected = chi2_contingency(observed_data)
            print(f"P-value ({quartile_1} and {quartile_2}): {p_value:.2f}")
```

P-value (Q1 and Q2): 0.47
P-value (Q1 and Q3): 0.23
P-value (Q1 and Q4): 0.05
P-value (Q2 and Q1): 0.47
P-value (Q2 and Q3): 0.72
P-value (Q2 and Q4): 0.25
P-value (Q3 and Q1): 0.23
P-value (Q3 and Q2): 0.72
P-value (Q3 and Q4): 0.52
P-value (Q4 and Q1): 0.05
P-value (Q4 and Q2): 0.25
P-value (Q4 and Q3): 0.52

```
In [145... daily_activity_sleep.groupby("METs_Quartile")["METs"].describe()
```

Out[145]: **count** **mean** **std** **min** **25%** **50%** **75%** **max**

METs_Quartile									
Q1	99.0	1.246579	0.091446	1.018056	1.189826	1.276111	1.321001	1.362014	
Q2	99.0	1.468878	0.053821	1.363611	1.430764	1.478125	1.515382	1.546875	
Q3	98.0	1.612353	0.040929	1.549722	1.577535	1.612500	1.644080	1.685556	
Q4	97.0	1.854094	0.187586	1.688542	1.732083	1.777569	1.905486	2.501667	

The MET quartiles were classified as 'Lowest Activity' (1.018056 - 1.364236 average daily MET/minute), 'Light Activity' (1.369375 - 1.546875 average daily MET/minute), 'Moderate Activity' (1.549722 - 1.688889 average daily MET/minute), 'Highest Activity' (1.688542 - 2.501667 average daily MET/minute) categories from the lowest to the highest. Sleep duration was grouped into 'undersleeping' (<6 hrs), 'normal' (6-8 hrs), and 'oversleeping' (>8 hrs). Although it seems that people with the lowest physical activity seem to undersleep less

than other groups, it appears to be due to their higher tendency to oversleep. Even light physical activity seems to positively affect the duration of sleep favouring sleep duration within the optimal range.

There is also a trend showing an increase in the fraction of sleep events with normal duration as physical activity increases. The difference between the fraction of normal duration sleep events in the lowest activity and highest activity quartiles did not pass the threshold of statistical significance. However, the p-value was very close to the threshold ($p = 0.06$), suggesting that the difference between the quartiles might be real. As the power of chi-square test increases with sample size, a higher sample size could be used to verify that the difference between the bottom and top quartiles is real.

```
daily_activity_sleep.to_csv("activity_sleep_data_cleaned.csv")
```

Average Hourly Intensities

```
In [146...] hourly_intensities["ActivityHour"] = hourly_intensities["ActivityHour"].dt.time
```

```
In [147...] hourly_intensities_grouped = hourly_intensities.groupby("ActivityHour").mean().reset_index()
hourly_intensities_grouped["ActivityHourPlot"] = hourly_intensities_grouped["ActivityHour"]
hourly_intensities_grouped.dtypes
```

```
Out[147]: ActivityHour      object
AverageIntensity    float64
ActivityHourPlot      int64
dtype: object
```

```
In [148...] hourly_intensities_grouped["ActivityHour"] = hourly_intensities_grouped["ActivityHourPlot"]
```


```
In [149...] x = hourly_intensities_grouped["ActivityHourPlot"]
y = hourly_intensities_grouped["AverageIntensity"]

plt.bar(x, y)

plt.xlabel("Time of a Day")
plt.ylabel("Average Intensity")

plt.xticks(hourly_intensities_grouped["ActivityHourPlot"][1::2], hourly_intensities_grouped["ActivityHourPlot"][1::2].dt.time)

plt.show()
```



hourly_intensities_grouped.to_csv("hourly_intensities_cleaned.csv")

Calculating the Level of Physical Activity Required to Achieve 1.5 Daily Average METs/minute

```
In [150...] average_daily_met["METs"].describe()
```

```
Out[150]: count    934.000000
mean        1.466396
std         0.290326
min         1.000000
25%         1.271024
50%         1.469618
75%         1.640677
max         2.577569
Name: METs, dtype: float64
```

```
In [151...] from sympy import *

for MET in [3, 6, 9]:
    x, y = symbols('x y')
    eq1 = Eq((1*x + MET*y)/24, 1.5)
    eq2 = Eq(x + y, 24)
    solution = solve((eq1,eq2), (x, y))
    print(f"{solution[y]:.2f} hours of activity at {MET} METs")
```

6.00 hours of activity at 3 METs

2.40 hours of activity at 6 METs

1.50 hours of activity at 9 METs

To reach 1.5 daily average METs/minute associated with more consistent sleep duration and more sleep events with a duration that falls within the normal range, people need to accumulate at least 6 hours of light activity, like slow walking, or 2.4 hours of moderate activity, like brisk walking, or 1.5 hours of intense physical activity, like running (<https://doi.org/10.1371%2Fjournal.pone.0200701>). This seems to be quite feasible as the average daily METs/minute of people who participated in this study is 1.47.